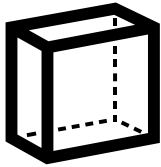


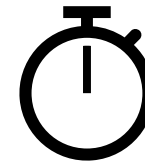
Coin Detector



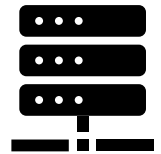
Technical Motivation for the Project



Object recognition as
an interesting field



Time intensive training
& prediction



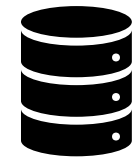
Computationally
intensive



Performing multiple
tasks

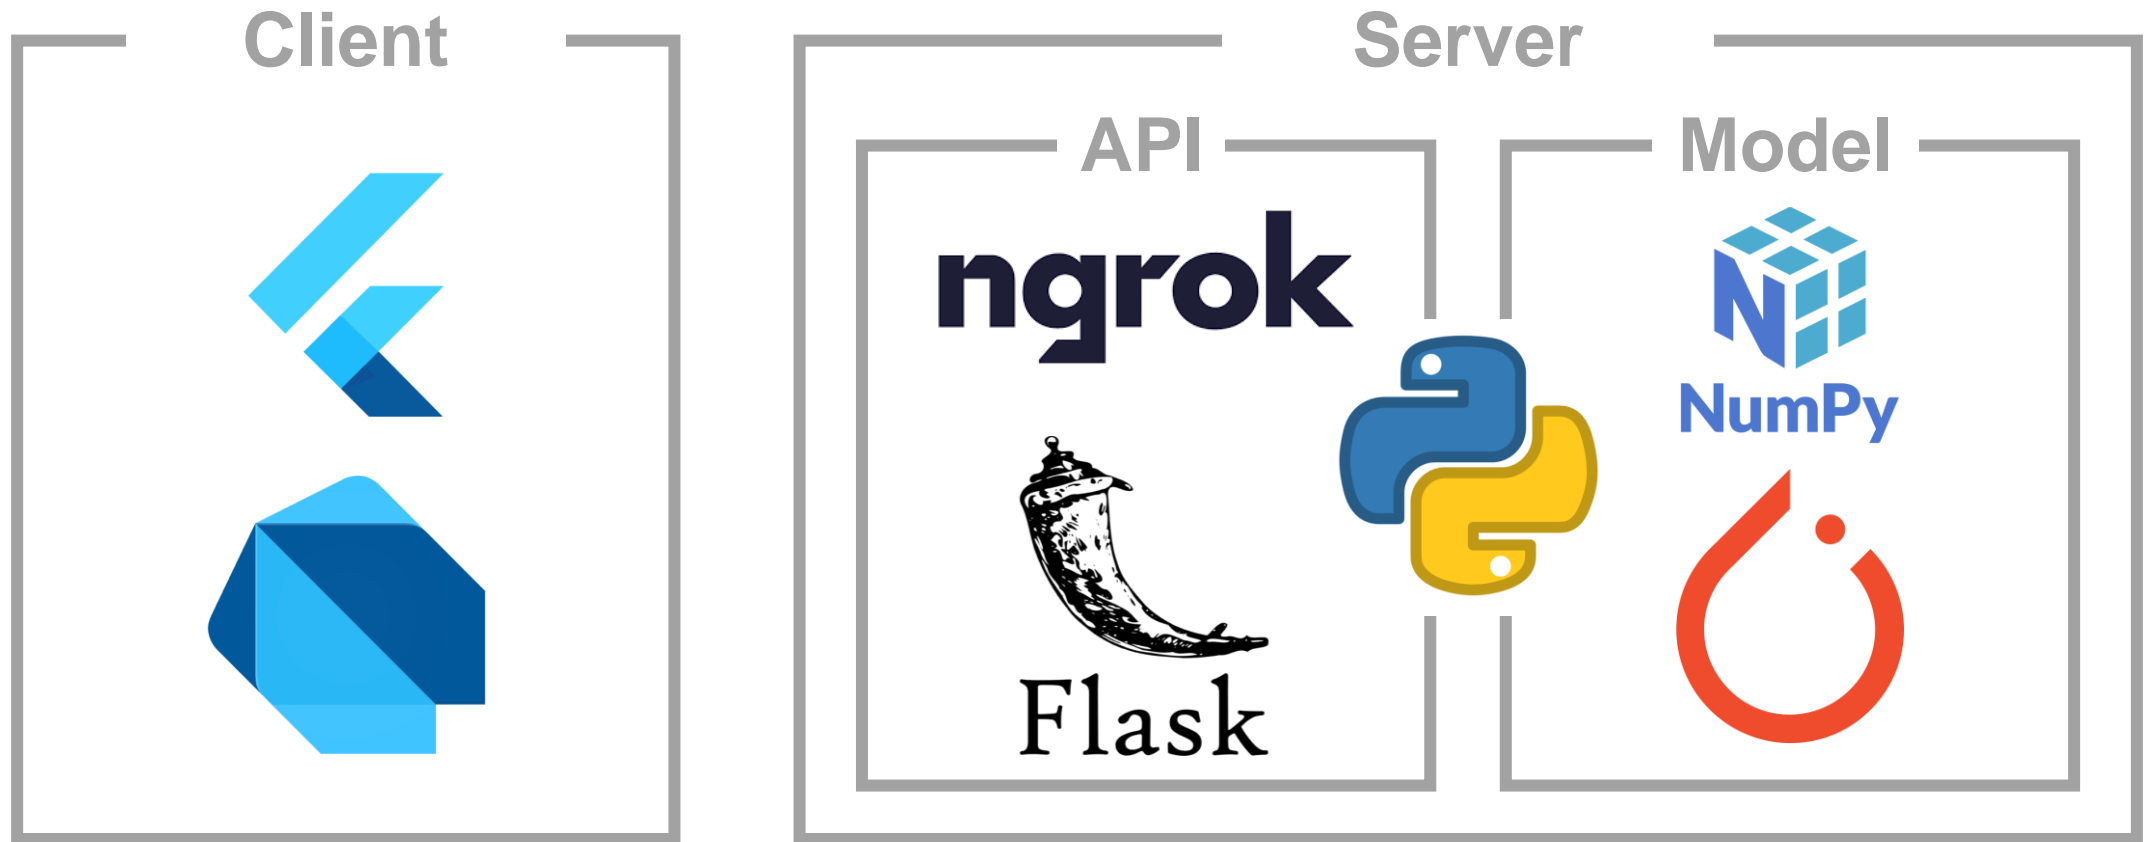


Communication
between client &
server

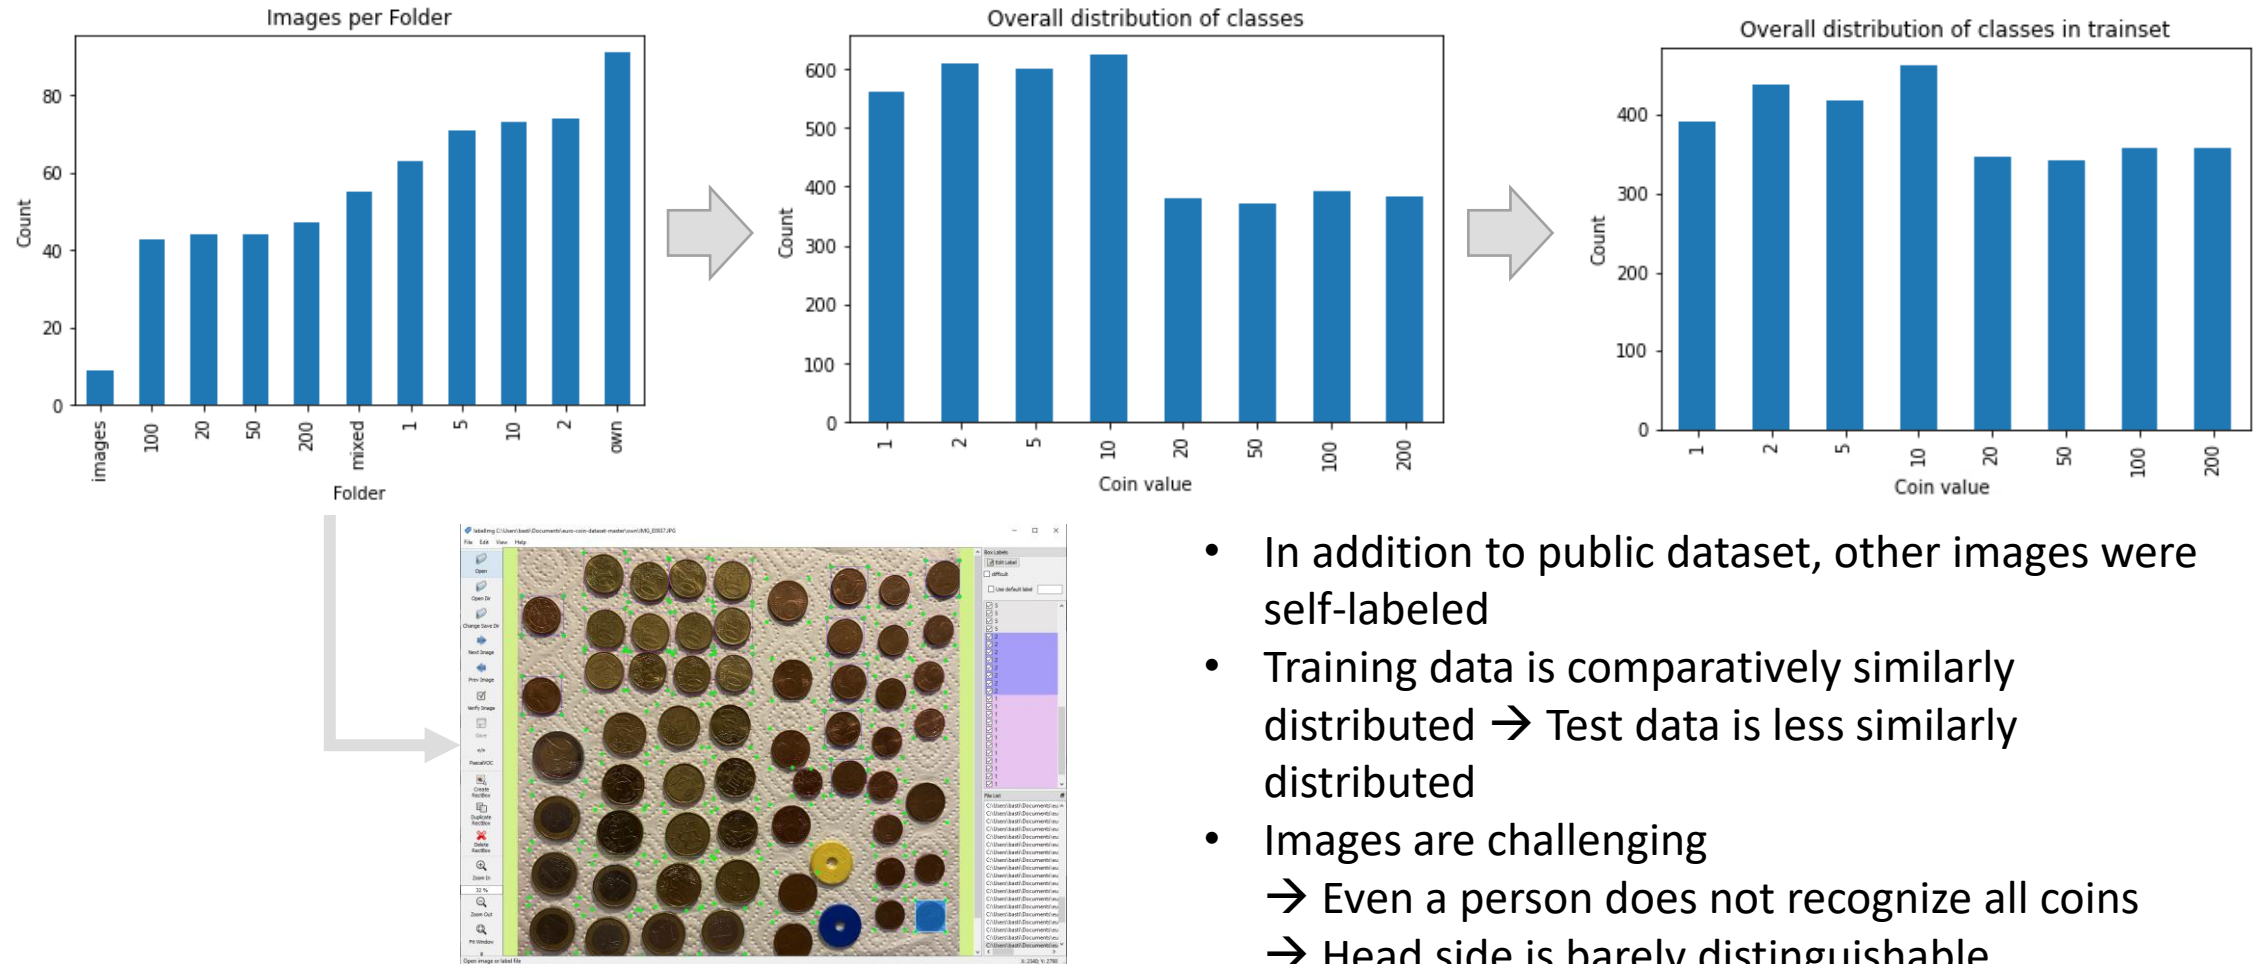


Data poses a challenge

Technology Stack



Data



- In addition to public dataset, other images were self-labeled
- Training data is comparatively similarly distributed → Test data is less similarly distributed
- Images are challenging
 - Even a person does not recognize all coins
 - Head side is barely distinguishable

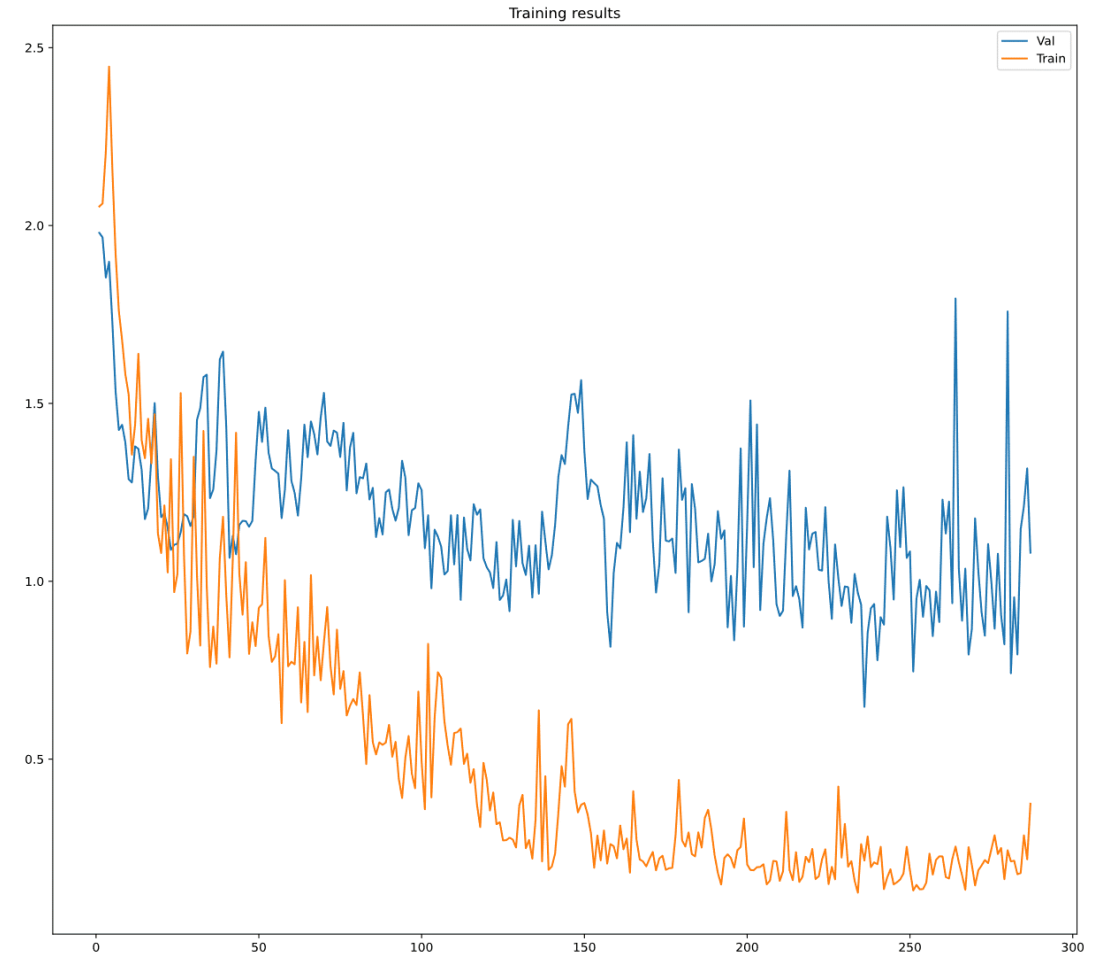
Model Selection & Training



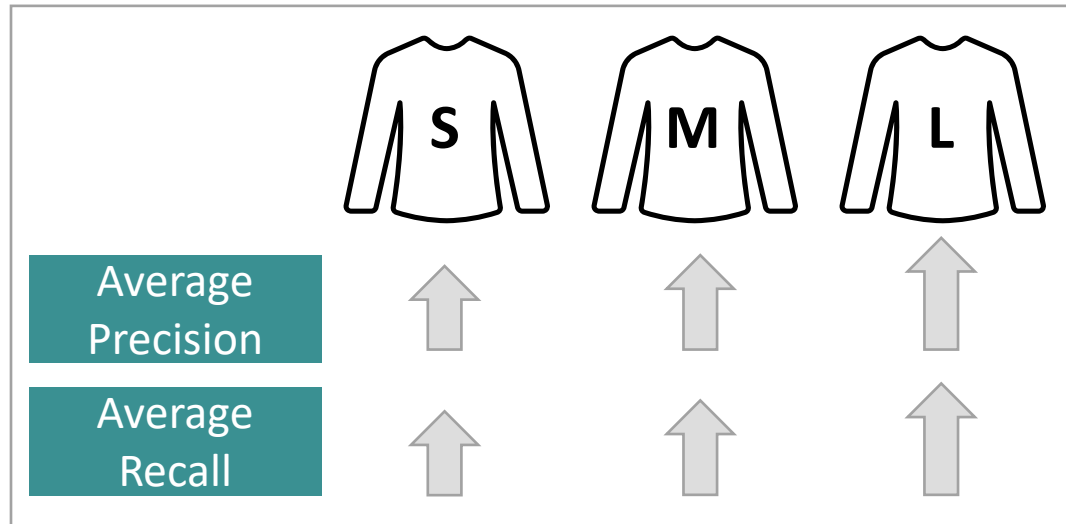
**Faster R-CNN
MobileNet v3**

PyTorch
Epochs: 287
Batch Size: 4
Pretrained: COCO
Graphics Card: P100

Fewer parameters than ResNet-50 or VGG-16, thus:
faster & less expensive training faster predictions
→ faster & cheaper training
→ faster predictions

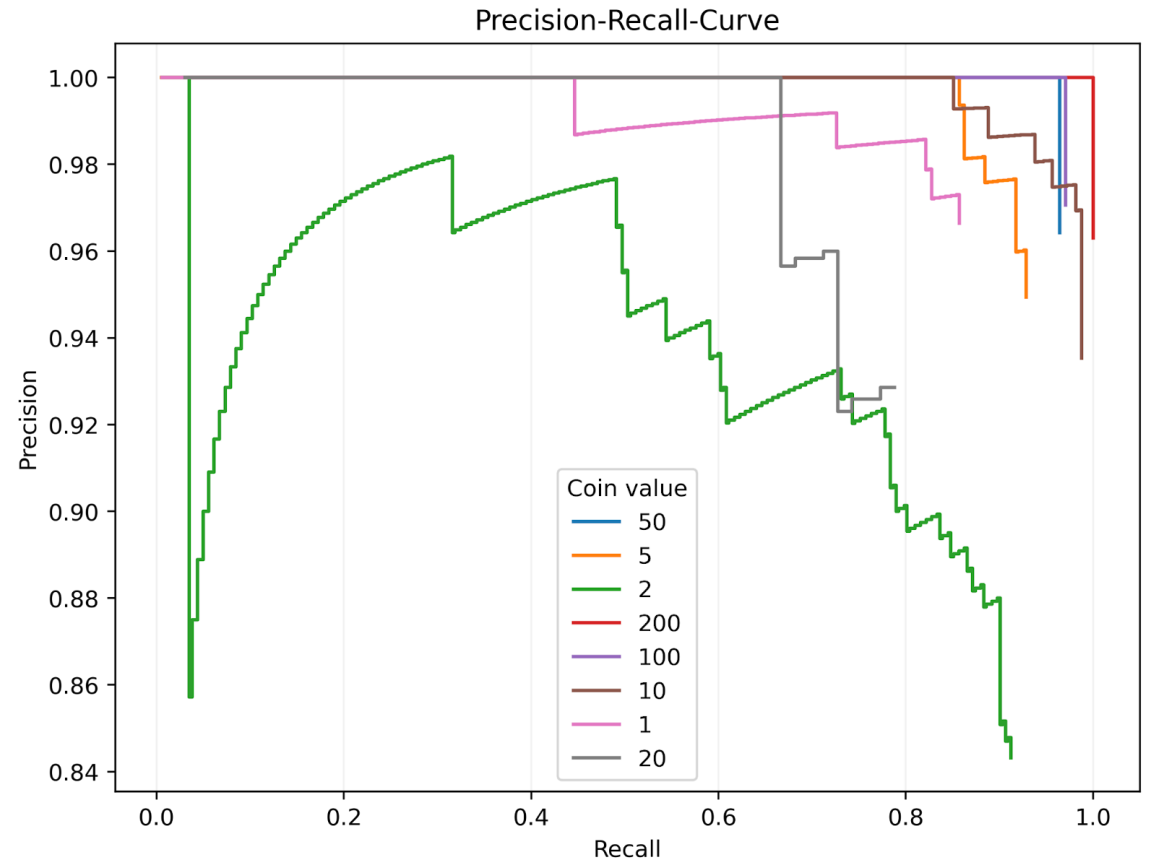


Model Evaluation

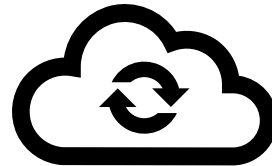
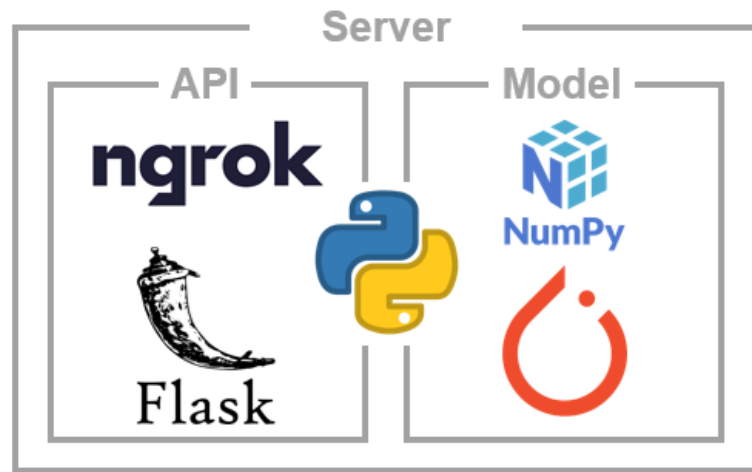


Findings:

- Average Precision of 0.79
- No large jumps between object sizes
- Big problems with 1 and 2 cent coins → Why?



Server

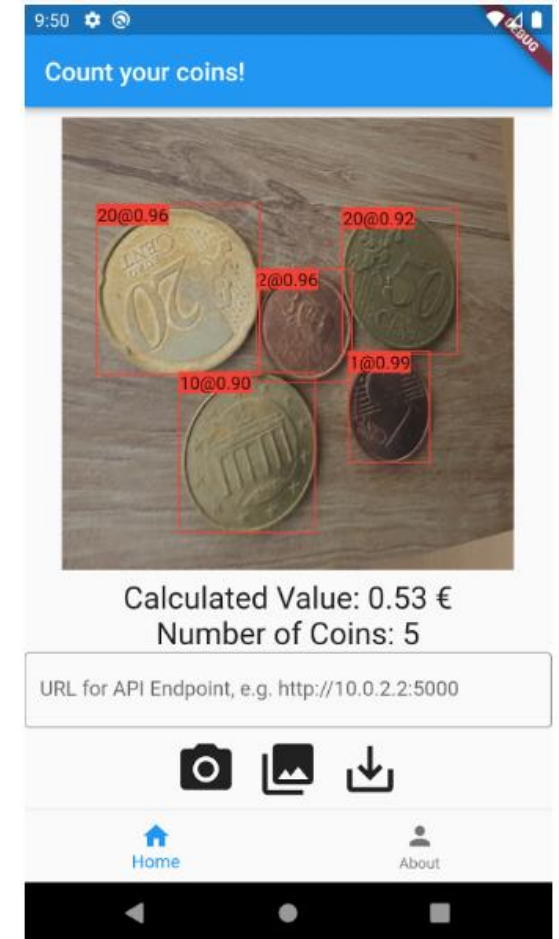


- The model is in the backend for several reasons:
 - Easy change
 - Performance → long inference time on cpu
 - Memory size of the model
 - Returned values :
 - Accumulated value of coins
 - Bounding boxes of recognized coins
 - Labels of recognized coins
 - Confidences for recognized coins
- Why not a fully rendered image?

App



- Take photo from gallery
- Take new photo
- Photos are allowed only in rectangular format
- Send image to backend
- Coordinates of bounding boxes, labels and confidences are sent back to the app
- Information displayed in app on image
- Save image with new info



Outlook

**Model deployment on
the client device**

**Host server on cloud
instance**

Label new images

iOS support

**Differentiation of the
training data in head or
number**

Publish app